

Merging Gentzkow Replication Data with Census Places

nwfried

2024-03-18

Data Wrangling

Initialising libraries.

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(fedmatch)
```

Loading in data.

```
cities <- read_tsv("30261-0006-Data.tsv")
```

```
## Rows: 2159 Columns: 4
## -- Column specification -----
## Delimiter: "\t"
## chr (2): cityname_constant, state
## dbl (2): citypermid, cnty90
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
census_places <- read_csv("places2msa1970.csv")
```

```
## Rows: 6584 Columns: 37
## -- Column specification -----
## Delimiter: ","
## chr (24): STATE, NHGISST, PLACE, NAME, NHGISPLACE, GISJOIN, NHGISNAM, NHGISS...
## dbl (12): YEAR, DECADE, ICPSRST, ICPSRCTY, ICPSRSTI, ICPSRCTYI, ICPSRFIP, PI...
## lgl (1): entityfips
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

First, we need to deal with the fact that Gentzkow's cities dataframe contains city names in all caps and the Census dataframe doesn't. This won't work with fuzzy match, so let's format both of these in all lowercase.

```

cities$cityname_constant <- tolower(cities$cityname_constant)
census_places$NAME <- tolower(census_places$NAME)
census_places$PLACE <- tolower(census_places$PLACE)

```

Next, we can fuzzy match based on the "NAME" column.

```

fuzzy_merge <- merge_plus(data1 = census_places,
  data2 = cities,
  by.x = "NAME",
  by.y = "cityname_constant", match_type = "fuzzy",
  unique_key_1 = "NHGISPLACE", unique_key_2 = "citypermid")

name_result <- fuzzy_merge$matches

```

Let's look at this result.

```
head(name_result)
```

```

## Key: <citypermid>
##   citypermid NHGISPLACE   STATE NHGISST   PLACE YEAR GISJOIN
##   <num>      <char>    <char> <char>    <char> <num> <char>
## 1:         1 G170407670 Illinois   170 la grange village 1970 G1703065
## 2:         1 G390412300  Ohio     390 lagrange village 1970 G3902085
## 3:         2 G390247660  Ohio     390 eldorado village 1970 G3901275
## 4:         4 G200200750 Kansas   200 el dorado city 1970 G2000780
## 5:         6 G020661400 Alaska    020 st. mary's city 1970 G0202860
## 6:         9 G100508000 Delaware  100 new castle city 1970 G1000200
##   DECADE NHGISNAM NHGISST_2 NHGISCTY ICPSRST ICPSRCTY ICPSRNAM STATENAM
##   <num>   <char>   <char> <char> <num> <num> <char> <char>
## 1:  1970    Cook    170   0310    21   310    COOK Illinois
## 2:  1970   Lorain   390   0930    24   930   LORAIN  Ohio
## 3:  1970   Greene   390   0570    24   570   GREENE  Ohio
## 4:  1970   Butler   200   0150    32   150   BUTLER  Kansas
## 5:  1970   Juneau   020   1100    81  1100   JUNEAU  Alaska
## 6:  1970 New Castle  100   0030    11   30 NEW CASTLE Delaware
##   ICPSRSTI ICPSRCTYI ICPSRFIP STATE_2 COUNTY  PID X_CENTROID Y_CENTROID
##   <num>     <num>     <num> <char> <char> <num> <num> <num>
## 1:     21        310         0   170   0310   596  673976.8  515200.60
## 2:     24        930         0   390   0930  1228 1146699.6  508982.44
## 3:     24        570         0   390   0570  1213 1026419.6  311105.56
## 4:     32        150         0   200   0150  1007  -73203.1  31813.45
## 5:     81       1100         0 <NA> <NA>     0     0.0     0.00
## 6:     11        30         0   100   0030   277 1718758.8  417391.14
##   GISJOIN_2 GISJOIN2 SHAPE_AREA SHAPE_LEN statefips countyfips fips smsacode
##   <char>    <char>    <num>    <num> <char>    <char> <char> <char>
## 1: G1700310 1700310 2481420600 303706.3    17     031 17031 1600
## 2: G3900930 3900930 1281184984 172594.8    39     093 39093 4440
## 3: G3900570 3900570 1077986175 137720.1    39     057 39057 2000
## 4: G2000150 2000150 3746230669 246761.8    20     015 20015 9040
## 5: G0201100 0201100 5028498768 1235933.0 <NA>    <NA> <NA> <NA>
## 6: G1000030 1000030 1121606309 220727.1    10     003 10003 9160
##   statefips_ countyfi_1 entityfips   name_2
##   <char>    <char>    <lgcl>    <char>
## 1:     17        031      NA      Cook County
## 2:     39        093      NA      Lorain County

```

```
## 3:      39      057      NA      Greene County
## 4:      20      015      NA      Butler County
## 5:     <NA>     <NA>      NA      <NA>
## 6:      10      003      NA New Castle County, DE
##
##              namemsa fips_2      NAME cityname_constant cnty90
##              <char> <char>      <char>      <char> <num>
## 1:           Chicago, IL SMSA  17031  la grange          lagrange 13285
## 2:           Lorain-Elyria, OH SMSA  39093  lagrange          lagrange 13285
## 3:           Dayton, OH SMSA  39057  eldorado          eldorado 17165
## 4:           Wichita, KS SMSA  20015  el dorado          el dorado 20015
## 5: Wilkes-Barre--Hazleton, PA SMSA  <NA> st. mary's      st. mary's 39011
## 6:           Wilmington, DE-NJ-MD SMSA  10003 new castle      new castle 42073
##
##      state tier
##      <char> <char>
## 1:    GEORGIA  all
## 2:    GEORGIA  all
## 3:   ILLINOIS  all
## 4:    KANSAS  all
## 5:     OHIO  all
## 6: PENNSYLVANIA  all
```

The problem here is that many place names are often reused, so the fuzzy match will match cities with the same name in different states e.g. Springfield, Illinois is matched with all other Springfields. We can use the county numbers of each city (found in the `cnty90` and `fips_2` columns, respectively) to make sure that the cities matched by the fuzzy match are actually the same. (we could also check this with state names, but I thought more specificity would be better at avoiding false positives).

First, we have to account for the fact that the four-digit county numbers in the Census dataframe are prefixed with a “0,” while the ones in the city database are not. Let’s add this 0 to the city database.

```
cnty_number <- cities$cnty90
fixed_cnty_number <- sprintf("%05d", cnty_number) #fix to 5 character width
cities <- cbind(fixed_cnty_number, cities) #add to dataset
```

Now we can do another fuzzy match and then filter out only the entries with matching county numbers.

```
merge2 <- merge_plus(data1 = census_places,
  data2 = cities,
  by.x = "NAME",
  by.y = "cityname_constant", match_type = "fuzzy",
  unique_key_1 = "NHGISPLACE", unique_key_2 = "citypermid")
result2 <- merge2$matches
unmatched_census <- merge2$data1_nomatch
unmatched_cities <- merge2$data2_nomatch

filtered_result2 <- filter(result2, fixed_cnty_number == fips_2)
head(filtered_result2)
```

```
## Key: <citypermid>
##   citypermid NHGISPLACE      STATE NHGISST      PLACE  YEAR
##   <num>      <char>      <char> <char>      <char> <num>
## 1:         4 G200200750    Kansas    200    el dorado city 1970
## 2:        39 G060165320  California    060    costa mesa city 1970
## 3:        40 G060690000  California    060    santa ana city 1970
## 4:        57 G120240000    Florida    120 fort lauderdale city 1970
## 5:       108 G250633050 Massachusetts    250    southbridge cdp 1970
## 6:       109 G250760300 Massachusetts    250    westfield city 1970
```

```

##      GISJOIN DECADE  NHGISNAM NHGISST_2 NHGISCTY ICPSRST ICPSRCTY  ICPSRNM
##      <char>  <num>    <char>    <char>    <char>    <num>    <num>    <char>
## 1: G2000780  1970    Butler      200      0150      32      150    BUTLER
## 2: G0600625  1970    Orange       060      0590      71      590    ORANGE
## 3: G0602570  1970    Orange       060      0590      71      590    ORANGE
## 4: G1200645  1970    Broward      120      0110      43      110    BROWARD
## 5: G2503980  1970    Worcester    250      0270       3      270    WORCESTER
## 6: G2504690  1970    Hampden     250      0130       3      130    HAMPDEN
##      STATENAM ICPSRSTI ICPSRCTYI ICPSRFIP STATE_2 COUNTY  PID X_CENTROID
##      <char>    <num>    <num>    <num>    <char> <char> <num>    <num>
## 1:      Kansas      32      150        0      200  0150  1007  -73203.1
## 2:    California      71      590        0      060  0590  2907 -1985465.2
## 3:    California      71      590        0      060  0590  2907 -1985465.2
## 4:      Florida      43      110        0      120  0110   285  1557532.5
## 5: Massachusetts      3      270        0      250  0270  2769  1950999.8
## 6: Massachusetts      3      130        0      250  0130  2762  1899443.2
##      Y_CENTROID GISJOIN_2 GISJOIN2 SHAPE_AREA SHAPE_LEN statefips countyfips
##      <num>    <char>    <char>    <num>    <num>    <char>    <char>
## 1:    31813.45 G2000150  2000150  3746230669  246761.8        20        015
## 2:   -196678.07 G0600590  0600590  2067710610  234799.8         06        059
## 3:   -196678.07 G0600590  0600590  2067710610  234799.8         06        059
## 4:  -1134960.10 G1200110  1200110  3162680337  247973.3         12        011
## 5:    792018.33 G2500270  2500270  4089604891  356649.8         25        027
## 6:    753974.26 G2500130  2500130  1642345438  274455.1         25        013
##      fips smsacode statefips_ countyfi_1 entityfips      name_2
##      <char>  <char>    <char>    <char>    <lgcl>      <char>
## 1:  20015      9040        20        015        NA      Butler County
## 2:  06059      0360         6         059        NA      Orange County
## 3:  06059      0360         6         059        NA      Orange County
## 4:  12011      2680        12        011        NA      Broward County
## 5:  25027      2600        25        027        NA      Worcester County (pt.)
## 6:  25013      8000        25        013        NA      Hampden County, MA (pt.)
##
##      namemsa fips_2      NAME
##      <char> <char>    <char>
## 1:      Wichita, KS SMSA  20015      el dorado
## 2:  Anaheim-Santa Ana-Garden Grove, CA SMSA  06059      costa mesa
## 3:  Anaheim-Santa Ana-Garden Grove, CA SMSA  06059      santa ana
## 4:      Fort Lauderdale-Hollywood, FL SMSA  12011      fort lauderdale
## 5:      Fitchburg-Leominster, MA SMSA  25027      southbridge
## 6: Springfield-Chicopee-Holyoke, MA-CT SMSA  25013      westfield
##      cityname_constant fixed_cnty_number cnty90      state  tier
##      <char>    <char>    <num>    <char> <char>
## 1:      el dorado      20015  20015      KANSAS  all
## 2:      costa mesa      06059  6059      CALIFORNIA  all
## 3:      santa ana      06059  6059      CALIFORNIA  all
## 4:      fort lauderdale      12011  12011      FLORIDA  all
## 5:      southbridge      25027  25027      MASSACHUSETTS  all
## 6:      westfield      25013  25013      MASSACHUSETTS  all

```

This gives us a list of 289 matched cities and census places. This is around a tenth of size of the original Gentzkow dataset and a fifth of the size of the initial fuzzy match merge. I'm wondering if there's some issue with the county filtering in that it's a bit too strict, given that ~1000 matches were dropped because the county numbers didn't match.

We can try filtering by state to see if this relaxes restrictions a bit. Again, we must set both datasets's state

column to lowercase to deal with case sensitivity.

```
cities$state <- tolower(cities$state)
census_places$STATE <- tolower(census_places$STATE)

state_merge <- merge_plus(data1 = census_places,
  data2 = cities,
  by.x = "NAME",
  by.y = "cityname_constant", match_type = "fuzzy",
  unique_key_1 = "NHGISPLACE", unique_key_2 = "citypermid") # merge again with changes to data
state_result <- state_merge$matches

filtered_state_result <- filter(state_result, state == STATE) # create dataset with matching states
head(filtered_state_result)
```

```
## Key: <citypermid>
##   citypermid NHGISPLACE      STATE NHGISST      PLACE  YEAR
##   <num>      <char>      <char> <char>      <char> <num>
## 1:         4 G200200750    kansas    200    el dorado city 1970
## 2:        11 G420624160 pennsylvania 420    pottstown borough 1970
## 3:        39 G060165320    california 060    costa mesa city 1970
## 4:        40 G060690000    california 060    santa ana city 1970
## 5:        57 G120240000    florida    120    fort lauderdale city 1970
## 6:       108 G250633050 massachusetts 250    southbridge cdp 1970
##   GISJOIN DECADE  NHGISNAM NHGISST_2 NHGISCTY ICPSRST ICPSRCTY ICPSRNAM
##   <char> <num>    <char>    <char>    <char> <num>    <num>    <char>
## 1: G2000780  1970    Butler    200    0150    32      150    BUTLER
## 2: G4207342  1970    Chester    420    0290    14      290    CHESTER
## 3: G0600625  1970    Orange    060    0590    71      590    ORANGE
## 4: G0602570  1970    Orange    060    0590    71      590    ORANGE
## 5: G1200645  1970    Broward    120    0110    43      110    BROWARD
## 6: G2503980  1970    Worcester    250    0270    3       270    WORCESTER
##   STATENAM ICPSRSTI ICPSRCTYI ICPSRFIP STATE_2 COUNTY  PID X_CENTROID
##   <char>    <num>    <num>    <num>    <char> <char> <num>    <num>
## 1: Kansas    32      150      0      200    0150  1007    -73203.1
## 2: Pennsylvania 14      290      0      420    0290  1477    1701418.1
## 3: California 71      590      0      060    0590  2907    -1985465.2
## 4: California 71      590      0      060    0590  2907    -1985465.2
## 5: Florida   43      110      0      120    0110   285    1557532.5
## 6: Massachusetts 3       270      0      250    0270  2769    1950999.8
##   Y_CENTROID GISJOIN_2 GISJOIN2 SHAPE_AREA SHAPE_LEN statefips countyfips
##   <num>      <char>    <char>    <num>    <num>    <char>    <char>
## 1: 31813.45 G2000150  2000150 3746230669 246761.8    20      015
## 2: 459142.37 G4200290  4200290 1967886858 250279.4    42      029
## 3: -196678.07 G0600590  0600590 2067710610 234799.8    06      059
## 4: -196678.07 G0600590  0600590 2067710610 234799.8    06      059
## 5: -1134960.10 G1200110  1200110 3162680337 247973.3    12      011
## 6: 792018.33 G2500270  2500270 4089604891 356649.8    25      027
##   fips smsacode statefips_ countyfi_1 entityfips      name_2
##   <char> <char>    <char>    <char>    <lgcl>      <char>
## 1: 20015    9040      20      015      NA      Butler County
## 2: 42029    6160      42      029      NA      Chester County, PA
## 3: 06059    0360      06      059      NA      Orange County
## 4: 06059    0360      06      059      NA      Orange County
## 5: 12011    2680      12      011      NA      Broward County
```

```
## 6: 25027      2600      25      027      NA Worcester County (pt.)
##              namemsa fips_2      NAME
##              <char> <char>      <char>
## 1:              Wichita, KS SMSA 20015      el dorado
## 2:              Philadelphia, PA-NJ SMSA 42029      pottstown
## 3: Anaheim-Santa Ana-Garden Grove, CA SMSA 06059      costa mesa
## 4: Anaheim-Santa Ana-Garden Grove, CA SMSA 06059      santa ana
## 5:      Fort Lauderdale-Hollywood, FL SMSA 12011 fort lauderdale
## 6:      Fitchburg-Leominster, MA SMSA 25027      southbridge
##      cityname_constant fixed_cnty_number cnty90      state      tier
##      <char>      <char> <num>      <char> <char>
## 1:      el dorado      20015 20015      kansas      all
## 2:      pottstown      42091 42091      pennsylvania      all
## 3:      costa mesa      06059 6059      california      all
## 4:      santa ana      06059 6059      california      all
## 5:      fort lauderdale      12011 12011      florida      all
## 6:      southbridge      25027 25027      massachusetts      all
```

```
state_merge_residue <- filter(state_result, state != STATE) # preserve data without matching states for
```

This returns around 540 matched cities and census places, filtering by state name.

Analysis

False matches

We can make it a little easier to look at this dataset by extracting only the most relevant columns.

```
extracted_state_result <- filtered_state_result[,c("STATE", "PLACE", "namemsa", "NAME", "cityname_constant")]
filter(extracted_state_result, cityname_constant != NAME)
```

```
##      STATE      PLACE      namemsa      NAME
##      <char>      <char>      <char>      <char>
## 1:      indiana      edinburgh town      Indianapolis, IN SMSA      edinburgh
## 2:      california      los altos city      San Jose, CA SMSA      los altos
## 3:      missouri      lee's summit city      Kansas City, MO-KS SMSA      lee's summit
##      cityname_constant      state
##      <char>      <char>
## 1:      edinburg      indiana
## 2:      los gatos      california
## 3:      lees summit      missouri
```

Filtering this dataset for place names that do not match yields only these three results. This should mean that the only entries with different city names are these three, of which the Los Altos/Los Gatos match seems to be the only false association.

Thus the only possible remaining false matches will be cities with the same name in the same state, which is hopefully only a handful. (I tried doing some research on this but I could only find various Reddit/Quora threads; there wasn't a definitive answer on how many of these places exist.)

Unmatched cities

This is where it gets a bit more tricky. First, we can look at the places that weren't matched at all in the fuzzy match: these seem to roughly fall into two categories.

- Places like North Attleboro, MA – there are a few cities near this in the unmatched Census data (same county) but no exact match. If we want to match these places, we could probably do this by looking at county number, but this might be dubious because they are ultimately different cities.
- Places like Dubois, PA – there’s nothing in this county in the Census data: presumably just not in the dataset.

There also seem to be some larger cities that are not getting matched because of differences in which they are labelled: for instance, “south chicago” exists in the unmatched_cities dataset and “south chicago heights” exists in the unmatched_census dataset. It might be prudent to go through some of the major cities to see if there are any other examples like this.

In general, major cities seem to be the biggest issue with this system of merging. I tried to look for New York to see how it was merged in this process, but I couldn’t find the original cities entry for “manhattan” in either the unmatched cities dataset or the matched dataset. This led me to believe that it had been incorrectly matched (say, with Manhattan, Kansas) and subsequently filtered out, so I created a dataset (state_merge_residue) to keep track of the cities that were matched and then filtered out because the state columns were different. But I couldn’t find the “manhattan” entry here either. I’m not sure where exactly this particular entry got removed, but it’s definitely possible that something similar has happened to other major cities.

Manual Matching

We can try to bring the major cities that have been dropped back into this merge manually.

Creating a dataframe to store our manual matches:

```
manual_matches <- head(cities, 0)
manual_matches <- cbind(manual_matches, head(census_places, 0))
# colnames(manual_matches) <- c(paste0("cities_", colnames(cities)), paste0("census_places_", colnames(census_places)))
colnames(manual_matches) <- c(paste0(colnames(cities)), paste0(colnames(census_places)))
```

Let’s first try matching Manhattan-New York manually:

```
# Manhattan
match_nyc_city <- cities[cities$citypermid == 148, ]
match_nyc_census <- census_places[census_places$PLACE == "new york city", ]
manual_matches <- rbind(manual_matches, c(match_nyc_city, match_nyc_census))
```

This is kind of tedious, we can write a function to speed it up a bit:

```
new_manual_match <- function(manual_matches, id_cities, id_census_places) {
  match_cities <- cities[cities$citypermid == id_cities, ]
  match_census <- census_places[census_places$PLACE == id_census_places, ]
  new_match <- c(match_cities, match_census)
  manual_matches <- rbind(manual_matches, new_match)
  return(manual_matches)
}
```

Let’s test the function:

```
# Test with Bronx-New York match
manual_matches <- new_manual_match(manual_matches, 1944, "new york city")
```

Now we can use a hashmap to store all of our manual matches and just pass that to the function. That way if we want to add a new match we can just add it to the hashmap.

```
manual_matches_map <- list(
  "queens" = list(permid = 2042, place= "new york city"),
  "brooklyn" = list(permid = 394, place= "new york city"),
```

```

"staten island" = list(permid = 1647, place = "new york city"),
"south chicago heights" = list(permid = 766, place = "south chicago heights village"),
"austin texas" = list(permid = 549, place = "austin city"),
"columbus ohio" = list(permid = 386, place = "columbus city"),
"indianapolis" = list(permid = 319, place = "indianapolis city (remainder)", # what does remainder m
"nashville" = list(permid = 47037, place = "nashville-davidson metropolitan government (balance)",
"las vegas" = list(permid = 1839, place = "las vegas city"),
"mesa" = list(permid = 20, place = "mesa city"),
"miami" = list(permid = 62, place = "miami city"), #census data contains both miami and miami beach
"virginia beach" = list(permid = 215, place = "virginia beach city") #include norfolk as well?
)

```

Iterate over this hashmap and add each entry to our manual_matches dataframe.

```

for (i in manual_matches_map) {
  id_cities <- i$permid
  id_census_places <- i$place
  match_cities <- cities[cities$citypermid == id_cities, ]
  match_census <- census_places[census_places$PLACE == id_census_places, ]
  new_match <- c(match_cities, match_census)
  manual_matches <- rbind(manual_matches, new_match)
}

```

Some places can't be identified by their census place name because there are multiple entries for one place name. Let's fix these manually:

```

manual_matches <- rbind(manual_matches, c(cities[cities$citypermid == 533, ], census_places[census_place
manual_matches <- rbind(manual_matches, c(cities[cities$citypermid == 2023, ], census_places[census_place

```

Now we have a dataframe complete with our manual matches that can easily be added to and appended to our main dataframe.